

# Design and Implementation of Anglican Probabilistic Programming Language

David Tolpin   Jan Willem van de Meent   Hongseok Yang  
Frank Wood

September 1, 2016

Paper and slides:

<https://bitbucket.org/probprog/anglican-white-paper>

# Outline

Why probabilistic programming?

Why functional?

# Intuition

## Probabilistic program:

- ▶ A program with random computations.
- ▶ Distributions are conditioned by ‘observations’.
- ▶ Values of certain expressions are ‘predicted’ — **the output**.

Can be written in any language (extended by `sample` and `observe`).

## Example: Model Selection

```
1  (let [;; Guessing a distribution
2      dist (sample (categorical
3                    [[normal 1] [gamma 1]
4                     [uniform-continuous 1]
5                     [uniform-discrete 1]]))
6      a (sample (gamma 1 1))
7      b (sample (gamma 1 1))
8      d (dist a b)]
9
10  ;; Observing samples from the distribution
11  (loop [data data]
12      (when (seq data)
13          (let [[x & data] data]
14              (observe d x))
15              (recur data)))
16
17  ;; Predicting a, b and the distribution
18  (predict :a a)
19  (predict :b b)
20  (predict :d d))
```

# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

- ▶ Initially,  $\mathcal{P}$  expects no arguments.

# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

- ▶ Initially,  $\mathcal{P}$  expects no arguments.
- ▶ On every call,  $\mathcal{P}$  returns
  - ▶ a distribution  $F$ ,
  - ▶ a distribution and a value  $(G, y)$ ,
  - ▶ a value  $z$ ,
  - ▶ or  $\perp$ .

# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

- ▶ Initially,  $\mathcal{P}$  expects no arguments.
- ▶ On every call,  $\mathcal{P}$  returns
  - ▶ a distribution  $F$ ,
  - ▶ a distribution and a value  $(G, y)$ ,
  - ▶ a value  $z$ ,
  - ▶ or  $\perp$ .
- ▶ Upon returning  $F$ ,  $\mathcal{P}$  expects  $x \sim F$ .



# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

- ▶ Initially,  $\mathcal{P}$  expects no arguments.
- ▶ On every call,  $\mathcal{P}$  returns
  - ▶ a distribution  $F$ ,
  - ▶ a distribution and a value  $(G, y)$ ,
  - ▶ a value  $z$ ,
  - ▶ or  $\perp$ .
- ▶ Upon returning  $F$ ,  $\mathcal{P}$  expects  $x \sim F$ .
- ▶ Upon returning  $\perp$ ,  $\mathcal{P}$  terminates.

# Definition

A **probabilistic program** is a stateful deterministic computation  $\mathcal{P}$ :

- ▶ Initially,  $\mathcal{P}$  expects no arguments.
- ▶ On every call,  $\mathcal{P}$  returns
  - ▶ a distribution  $F$ ,
  - ▶ a distribution and a value  $(G, y)$ ,
  - ▶ a value  $z$ ,
  - ▶ or  $\perp$ .
- ▶ Upon returning  $F$ ,  $\mathcal{P}$  expects  $x \sim F$ .
- ▶ Upon returning  $\perp$ ,  $\mathcal{P}$  terminates.

A program is run by calling  $\mathcal{P}$  repeatedly until termination.

The probability of each **trace** is  $\propto \prod_{i=1}^{|\mathbf{x}|} p_{F_i}(x_i) \prod_{j=1}^{|\mathbf{y}|} p_{G_j}(y_j)$ .

# Evolution of Compiler Development

- ▶ 1950s: Ask a programmer to implement an algorithm efficiently: They'll write it on their own in assembly.
  - ▶ No good compilers; problem-dependent optimizations that only human expert could see.

# Evolution of Compiler Development

- ▶ 1950s: Ask a programmer to implement an algorithm efficiently: They'll write it on their own in assembly.
  - ▶ No good compilers; problem-dependent optimizations that only human expert could see.
- ▶ 1970s: Novice programmers use high-level languages and let compiler work out details, experts still write assembly.
  - ▶ Experts still write custom assembly when speed critical.

# Evolution of Compiler Development

- ▶ 1950s: Ask a programmer to implement an algorithm efficiently: They'll write it on their own in assembly.
  - ▶ No good compilers; problem-dependent optimizations that only human expert could see.
- ▶ 1970s: Novice programmers use high-level languages and let compiler work out details, experts still write assembly.
  - ▶ Experts still write custom assembly when speed critical.
- ▶ 2000s: On most problems, even experts can't write faster assembly than optimizing compilers.
  - ▶ can automatically profile (JIT).
  - ▶ can take advantage of parallelization, complicated hardware, make appropriate choices w.r.t. caching.
  - ▶ Compilers embody decades of compiler research

# Evolution of Inference

- ▶ 2010: Ask a grad student to implement inference efficiently:  
They'll write it on their own.
  - ▶ No good automatic inference engines; problem-dependent optimizations that only human expert can see.

# Evolution of Inference

- ▶ 2010: Ask a grad student to implement inference efficiently: They'll write it on their own.
  - ▶ No good automatic inference engines; problem-dependent optimizations that only human expert can see.
- ▶ 2015: Novice grad students use automatic inference engines and let compiler work out details, experts still write their own inference.
  - ▶ Experts still write custom inference when speed critical.

# Evolution of Inference

- ▶ 2010: Ask a grad student to implement inference efficiently: They'll write it on their own.
  - ▶ No good automatic inference engines; problem-dependent optimizations that only human expert can see.
- ▶ 2015: Novice grad students use automatic inference engines and let compiler work out details, experts still write their own inference.
  - ▶ Experts still write custom inference when speed critical.
- ▶ 2020: On most problems, even experts can't write faster inference than mature automatic inference engines.
  - ▶ Can use parallelization, sophisticated hardware
  - ▶ Can automatically choose appropriate methods (meta-reasoning?).
  - ▶ Inference engines will embody 1 decade (!) of PP research.



# Outline

Why probabilistic programming?

Why functional?

# Inference Objective

- ▶ Suggest **most probable explanation** (MPE) - most likely assignment for all non-evidence variable given evidence.

# Inference Objective

- ▶ Suggest **most probable explanation** (MPE) - most likely assignment for all non-evidence variable given evidence.
- ▶ Approximately **compute integral** of the form

$$\Phi = \int_{-\infty}^{\infty} \varphi(x)p(x)dx$$




# Inference Objective

- ▶ Suggest **most probable explanation** (MPE) - most likely assignment for all non-evidence variable given evidence.
- ▶ Approximately **compute integral** of the form

$$\Phi = \int_{-\infty}^{\infty} \varphi(x)p(x)dx$$

- ▶ Continuously and **infinitely generate a sequence of samples** drawn from the distribution of the output expression — so that someone else puts it in good use (vague but common). ✓

# Example: Inference Results



# Importance Sampling

**loop**

Run  $\mathcal{P}$ , computing weight  $w = \prod_{j=1}^{|y|} p_{G_j}(y_j)$ .  
output  $\mathbf{z}$ ,  $w$ .

**end loop**

- ▶ Simple — good.
- ▶ Slow convergence (unless one knows the answer) — bad.

Can we do better?

# Lightweight Metropolis-Hastings (LMH)

Run  $\mathcal{P}$  once, remember  $\mathbf{x}, \mathbf{z}$ .

**loop**

Uniformly select  $x_i$ .

Propose a value for  $x_i$ .

Run  $\mathcal{P}$ , remember  $\mathbf{x}', \mathbf{z}'$ .

Accept  $(\mathbf{x}, \mathbf{z} = \mathbf{x}', \mathbf{z}')$  or reject with MH probability.

Output  $\mathbf{z}$ .

**end loop**

Can we do better?

- ▶ Particle Markov Chain Monte Carlo
- ▶ Variational Inference
- ▶ ...

# Challenges

- ▶ ‘Transformational compilation’ — limited languages, some less ugly than others.
- ▶ Slow inference — Markov Chain Monte Carlo is **always slow** (but there are Hybrid Monte Carlo, Variational Inference, ...)
- ▶ How to learn persistently — what is the learned model? Transfer learning?
- ▶ Handling indeterminism (learning policies).



Thank you!  
Questions?